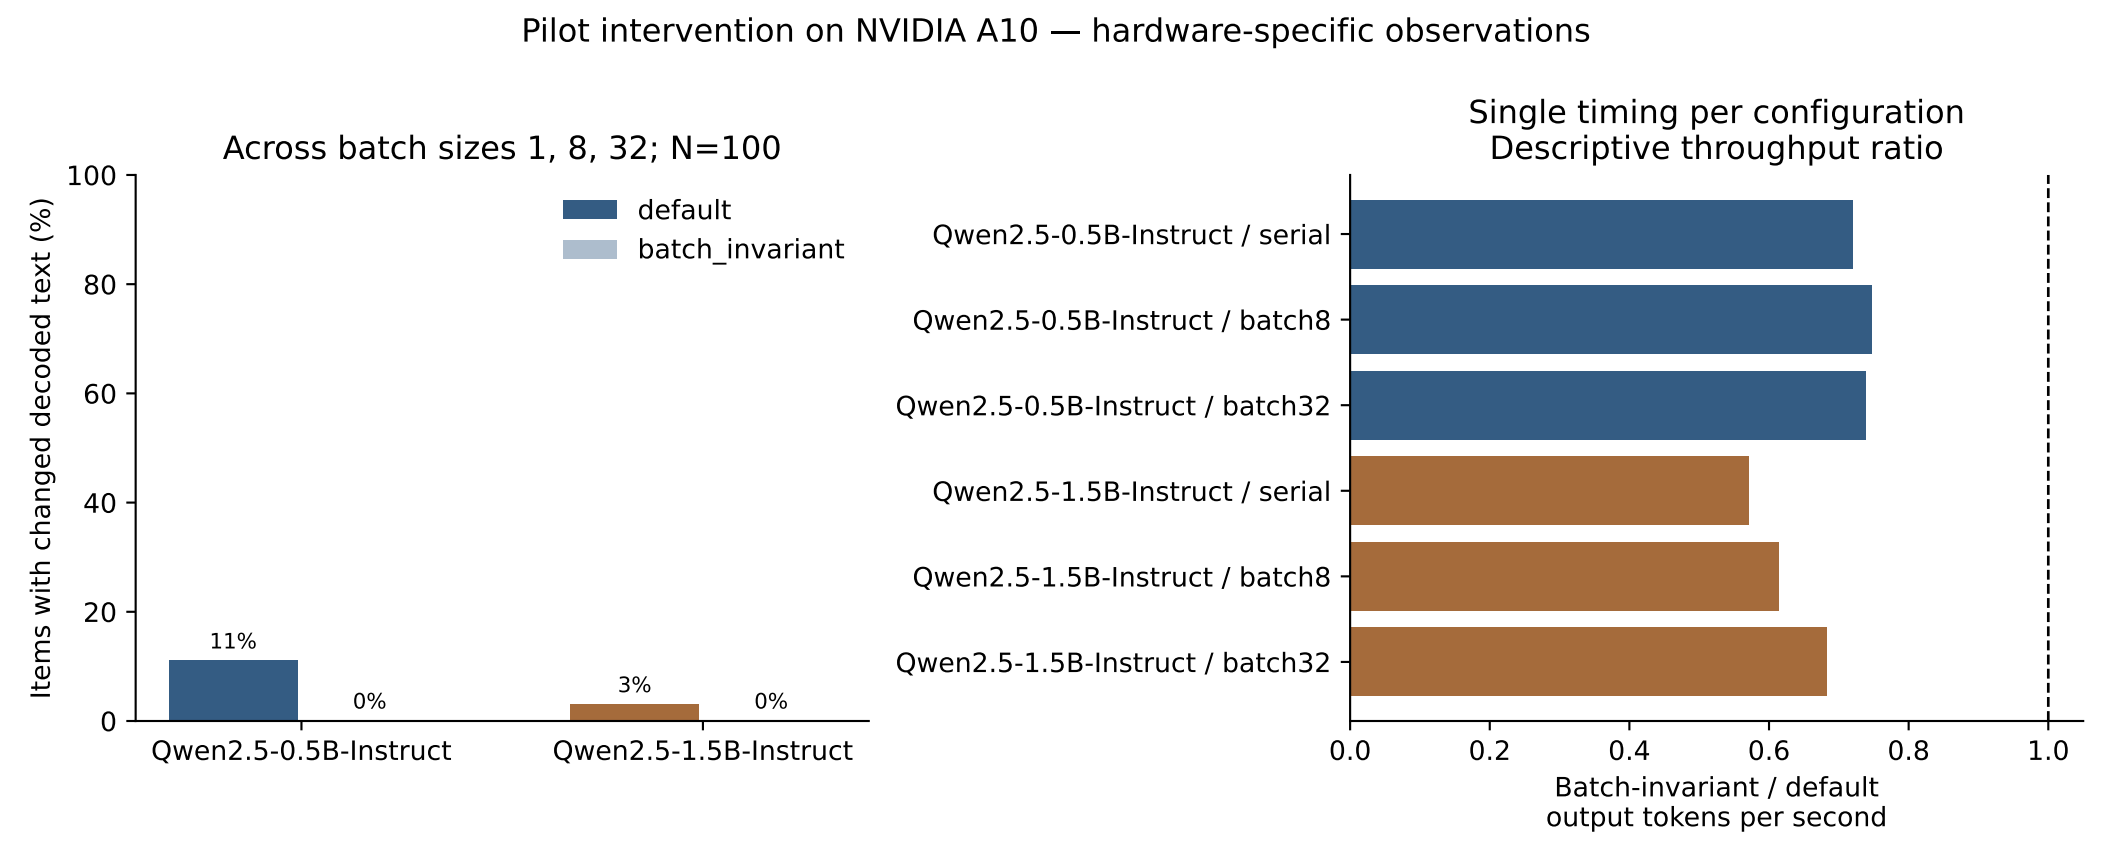
Pilot intervention on NVIDIA A10 — hardware-specific observations
Across batch sizes 1, 8, 32; N=100
Items with changed decoded text (%)
100
80
60
40
20
0
default
batch_invariant
11%
0%
3%
0%
Qwen2.5-0.5B-Instruct
Qwen2.5-1.5B-Instruct
Single timing per configuration
Descriptive throughput ratio
Qwen2.5-0.5B-Instruct / serial
Qwen2.5-0.5B-Instruct / batch8
Qwen2.5-0.5B-Instruct / batch32
Qwen2.5-1.5B-Instruct / serial
Qwen2.5-1.5B-Instruct / batch8
Qwen2.5-1.5B-Instruct / batch32
0.0
0.2
0.4
0.6
0.8
1.0
Batch-invariant / default
output tokens per second