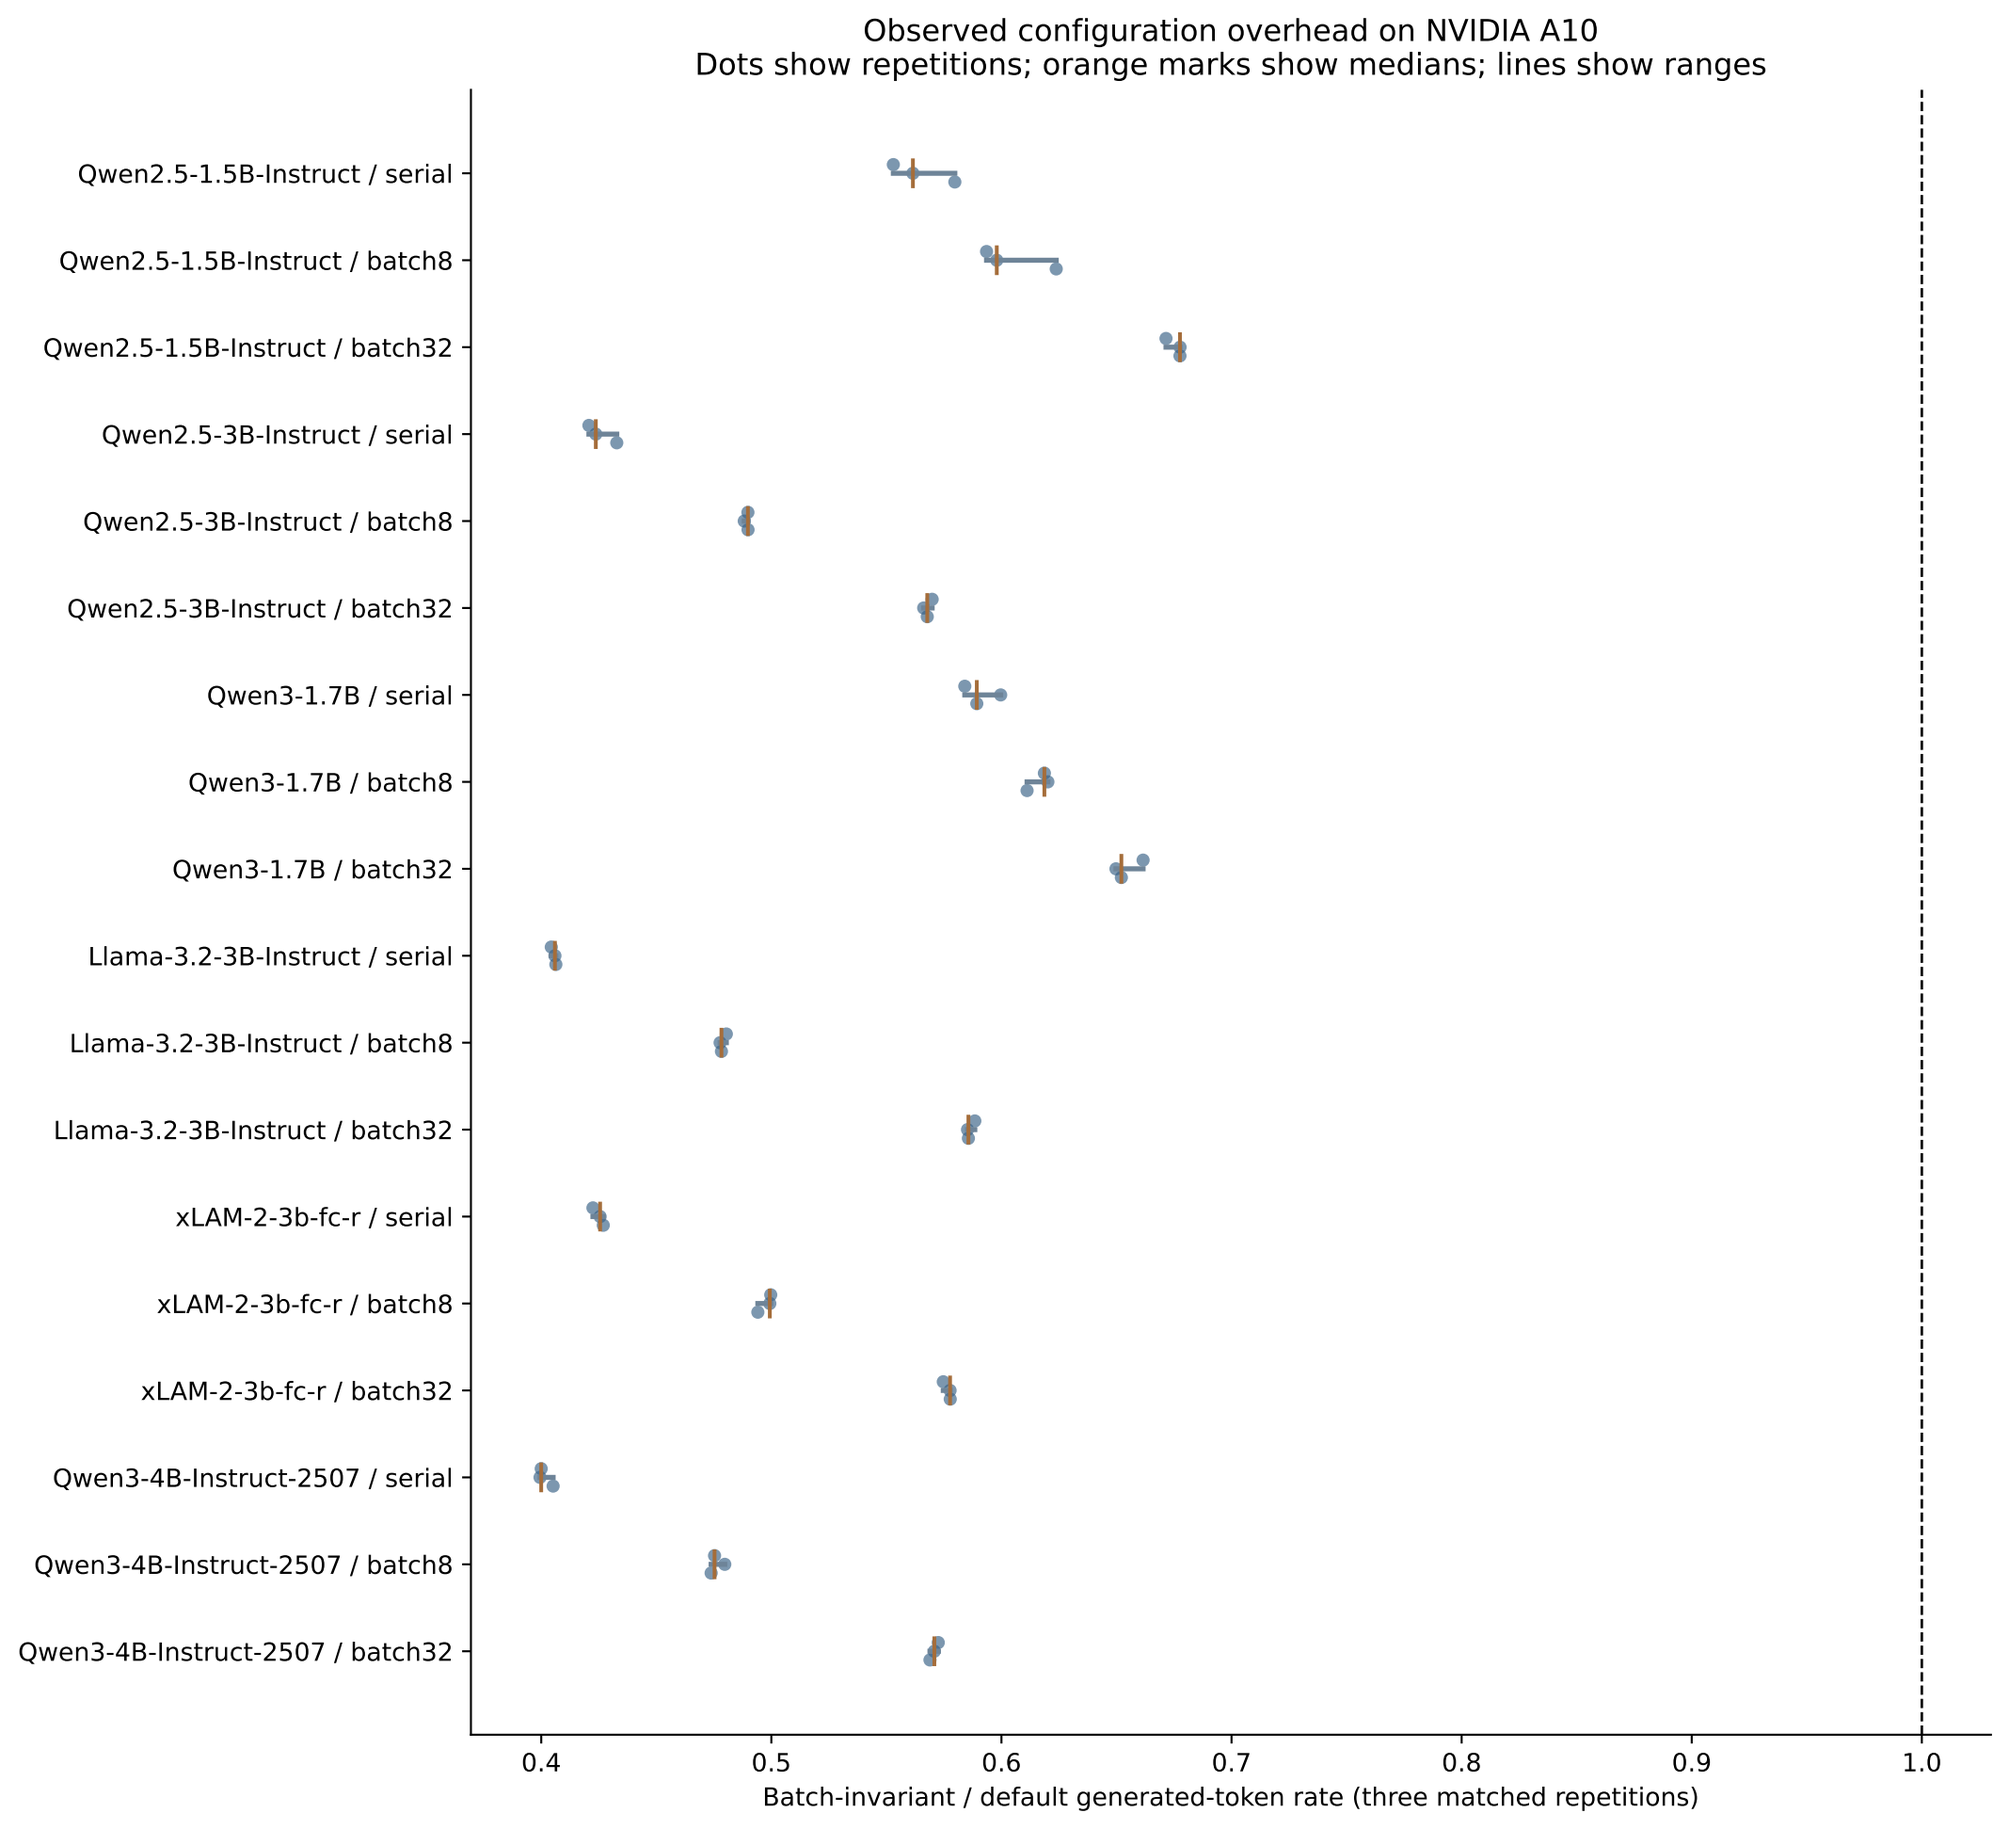
Observed configuration overhead on NVIDIA A10
Dots show repetitions; orange marks show medians; lines show ranges
Qwen2.5-1.5B-Instruct / serial
Qwen2.5-1.5B-Instruct / batch8
Qwen2.5-1.5B-Instruct / batch32
Qwen2.5-3B-Instruct / serial
Qwen2.5-3B-Instruct / batch8
Qwen2.5-3B-Instruct / batch32
Qwen3-1.7B / serial
Qwen3-1.7B / batch8
Qwen3-1.7B / batch32
Llama-3.2-3B-Instruct / serial
Llama-3.2-3B-Instruct / batch8
Llama-3.2-3B-Instruct / batch32
xLAM-2-3b-fc-r / serial
xLAM-2-3b-fc-r / batch8
xLAM-2-3b-fc-r / batch32
Qwen3-4B-Instruct-2507 / serial
Qwen3-4B-Instruct-2507 / batch8
Qwen3-4B-Instruct-2507 / batch32
0.4
0.5
0.6
0.7
0.8
0.9
1.0
Batch-invariant / default generated-token rate (three matched repetitions)